

## Graph Definitions

### Graphs

6/110

Many applications require

- a collection of *items* (i.e. a set)
- *relationships*/connections between items

Examples:

- maps: items are cities, connections are roads
- web: items are pages, connections are hyperlinks

Collection types you're familiar with

- arrays and lists ... linear sequence of items

*Graphs* are more general ... allow arbitrary connections


### ... Graphs

7/110

A graph  $G = (V, E)$

- $V$  is a set of *vertices*
- $E$  is a set of *edges* (subset of  $V \times V$ )

Example:



### ... Graphs

8/110

A real example: Australian road distances

Distance	Adelaide	Brisbane	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	-	2055	-	3051	732	2716	-
Brisbane	2055	-	-	3429	1671	-	982


Canberra	-	-	-	-	658	-	309
Darwin	3051	3429	-	-	-	4049	-
Melbourne	732	1671	658	-	-	-	873
Perth	2716	-	-	4049	-	-	-
Sydney	-	982	309	-	873	-	-

Notes: vertices are cities, edges are distance between cities, symmetric

## ... Graphs

9/110

Alternative representation of above:



## ... Graphs

10/110

Questions we might ask about a graph:

- is there a way to get from item A to item B?
- what is the best way to get from A to B?
- which items are connected?

Graph algorithms are generally more complex than tree/list ones:

- no implicit order of items
- graphs may contain cycles
- concrete representation is less obvious
- algorithm complexity depends on connection complexity

## Properties of Graphs

11/110

Terminology:  $|V|$  and  $|E|$  (cardinality) normally written just as  $V$  and  $E$ .

A graph with  $V$  vertices has at most  $V(V-1)/2$  edges.

The ratio  $E:V$  can vary considerably.

- if  $E$  is closer to  $V^2$ , the graph is *dense*

- if  $E$  is closer to  $V$ , the graph is *sparse*
  - Example: web pages and hyperlinks, intersections and roads on street map

Knowing whether a graph is sparse or dense is important

- may affect choice of data structures to represent graph
  - may affect choice of algorithms to process graph
- 

## Exercise #1: Number of Edges

12/110

The edges in a graph represent pairs of connected vertices. A graph with  $V$  has  $V^2$  such pairs.

Consider  $V = \{1,2,3,4,5\}$  with all possible pairs:

$E = \{ (1,1), (1,2), (1,3), (1,4), (1,5), (2,1), (2,2), \dots, (4,5), (5,5) \}$

Why do we say that the maximum #edges is  $V(V-1)/2$ ?

---

... because

- $(v,w)$  and  $(w,v)$  denote the same edge (in an undirected graph)
  - we do not consider loops  $(v,v)$  (in undirected graphs)
- 

## Graph Terminology

14/110

For an edge  $e$  that connects vertices  $v$  and  $w$

- $v$  and  $w$  are *adjacent* (neighbours)
- $e$  is *incident* on both  $v$  and  $w$

*Degree* of a vertex  $v$

- number of edges incident on  $v$

Synonyms:

- vertex = node, edge = arc = link (NB: some people use arc for *directed* edges)
- 

## ... Graph Terminology

15/110

*Path*: a sequence of vertices where

- each vertex has an edge to its predecessor

*Simple path*: a path where


- all vertices and edges are different

*Cycle*: a path

- that is simple except last vertex = first vertex

*Length* of path or cycle:

- #edges



## ... Graph Terminology

16/110

*Connected graph*

- there is a *path* from each vertex to every other vertex
- if a graph is not connected, it has  $\geq 2$  *connected components*

*Complete graph  $K_V$*

- there is an *edge* from each vertex to every other vertex
- in a complete graph,  $E = V(V-1)/2$



## ... Graph Terminology


17/110

*Tree*: connected (sub)graph with no cycles

*Spanning tree*: tree containing all vertices

*Clique*: complete subgraph

Consider the following single graph:



This graph has 26 vertices, 33 edges, and 4 connected components

Note: The entire graph has no spanning tree; what is shown in green is a spanning tree of the third connected component

## ... Graph Terminology

18/110

A *spanning tree* of connected graph  $G = (V, E)$

- is a subgraph of  $G$  containing all of  $V$
- and is a single tree (connected, no cycles)


A *spanning forest* of non-connected graph  $G = (V, E)$

- is a subgraph of  $G$  containing all of  $V$
- and is a set of trees (not connected, no cycles),
  - with one tree for each *connected component*

---

## Exercise #2: Graph Terminology

19/110



$$V = \{v1, v2, v3, v4\}$$

$$E = \{e1, e2, e3, e4, e5\}$$

1. How many edges need to be removed to obtain a spanning tree?
2. How many different spanning trees?

- 
1. 2
  2.  $\frac{5 \cdot 4}{2} - 2 = 8$  spanning trees (no spanning tree if we remove  $\{e1, e2\}$  or  $\{e3, e4\}$ )

---

## ... Graph Terminology

21/110


*Undirected graph*

- $edge(u, v) = edge(v, u)$ , no self-loops (i.e. no  $edge(v, v)$ )


*Directed graph*

- $edge(u, v) \neq edge(v, u)$ , can have self-loops (i.e.  $edge(v, v)$ )

Examples:



*Undirected graph*



*Directed graph*

---

## ... Graph Terminology

22/110

Other types of graphs ...

## Weighted graph

- each edge has an associated value (weight)
- e.g. road map (weights on edges are distances between cities)

## Multi-graph

- allow multiple edges between two vertices
- e.g. function call graph (`f()` calls `g()` in several places)

---

# Graph Data Structures

---


## Graph Representations

24/110

Defining graphs:

- need some way of identifying vertices
- could give diagram showing edges and vertices
- could give a list of edges

E.g. four representations of the same graph:



---

## ... Graph Representations

25/110

We will discuss three different graph data structures:

1. Array of edges
2. Adjacency matrix
3. Adjacency list


---

## Array-of-edges Representation


26/110

Edges are represented as an array of `Edge` values (= pairs of vertices)

- space efficient representation
- adding and deleting edges is slightly complex
- undirected: order of vertices in an `Edge` doesn't matter
- directed: order of vertices in an `Edge` encodes direction



[ (0,1), (1,2), (1,3), (2,3) ]



[ (1,0), (1,1), (0,2), (0,3), (2,3) ]

For simplicity, we always assume vertices to be numbered  $0..V-1$

---

## ... Array-of-edges Representation

27/110

### Graph initialisation

```

newGraph(V) :
|   Input  number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate enough memory for g.edges[]
|   return g
  
```

How much is enough? ... No more than  $V(V-1)/2$  ... Much less in practice (sparse graph)

---

## ... Array-of-edges Representation

28/110

### Edge insertion

```

insertEdge(g, (v,w)) :
|   Input graph g, edge (v,w)  // assumption: (v,w) not in g
|
|   g.edges[g.nE]=(v,w)
|   g.nE=g.nE+1
  
```

## ... Array-of-edges Representation

29/110

### Edge removal

```

removeEdge(g, (v,w)) :
|   Input graph g, edge (v,w)  // assumption: (v,w) in g
|
|   i=0
|   while (v,w) ≠ g.edges[i] do
|     i=i+1
|   end while
|   g.edges[i]=g.edges[g.nE-1] // replace (v,w) by last edge in array
|   g.nE=g.nE-1
  
```

---

## Cost Analysis

30/110

Storage cost:  $O(E)$

Cost of operations:

- initialisation:  $O(1)$
- insert edge:  $O(1)$  (assuming edge array has space)
- find/delete edge:  $O(E)$  (need to find edge in edge array)

If array is full on insert

- allocate space for a bigger array, copy edges across  $\Rightarrow O(E)$

If we maintain edges in order

- use binary search to insert/find edge  $\Rightarrow O(\log E)$

---

### Exercise #3: Array-of-edges Representation

31/110

Assuming an array-of-edges representation ...

Write an algorithm to output all edges of the graph

```
show(g):  
|   Input graph g  
|  
|   for all i=0 to g.nE-1 do  
|       print g.edges[i]  
|   end for
```


Time complexity:  $O(E)$

---

### Adjacency Matrix Representation


33/110

Edges represented by a  $V \times V$  matrix



Undirected graph

A	0	1	2	3
0	0	1	0	1
1	1	0	0	1
2	0	0	0	1
3	1	1	1	0



Directed graph

A	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	0	0	0	0
3	0	0	1	0

---

### ... Adjacency Matrix Representation

34/110

Advantages



- easily implemented as 2-dimensional array
- can represent graphs, digraphs and weighted graphs
  - graphs: symmetric boolean matrix
  - digraphs: non-symmetric boolean matrix
  - weighted: non-symmetric matrix of weight values

Disadvantages:

- if few edges (sparse)  $\Rightarrow$  memory-inefficient

## ... Adjacency Matrix Representation

35/110

### Graph initialisation

```
newGraph(V):
|   Input  number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate memory for g.edges[][]
|   for all i,j=0..V-1 do
|       g.edges[i][j]=0    // false
|   end for
|   return g
```

## ... Adjacency Matrix Representation

36/110

### Edge insertion

```
insertEdge(g, (v,w)):
|   Input graph g, edge (v,w)
|
|   if g.edges[v][w]=0 then // (v,w) not in graph
|       g.edges[v][w]=1    // set to true
|       g.edges[w][v]=1
|       g.nE=g.nE+1
|   end if
```

## ... Adjacency Matrix Representation

37/110

### Edge removal

```
removeEdge(g, (v,w)):
|   Input graph g, edge (v,w)
|
|   if g.edges[v][w]≠0 then // (v,w) in graph
|       g.edges[v][w]=0    // set to false
|       g.edges[w][v]=0
|       g.nE=g.nE-1
|   end if
```

## Exercise #4: Show Graph

38/110

Assuming an adjacency matrix representation ...

Write an algorithm to output all edges of the graph (no duplicates!)

### ... Adjacency Matrix Representation

39/110

```
show(g):  
    Input graph g  
  
    for all i=0 to g.nV-2 do  
        for all j=i+1 to g.nV-1 do  
            if g.edges[i][j] then  
                print i—"—"j  
            end if  
        end for  
    end for
```

Time complexity:  $O(V^2)$

## Exercise #5:

40/110

Analyse storage cost and time complexity of adjacency matrix representation

Storage cost:  $O(V^2)$

If the graph is sparse, most storage is wasted.


Cost of operations:

- initialisation:  $O(V^2)$  (initialise  $V \times V$  matrix)
- insert edge:  $O(1)$  (set two cells in matrix)
- delete edge:  $O(1)$  (unset two cells in matrix)

### ... Adjacency Matrix Representation

42/110

A storage optimisation: store only top-right part of matrix.




New storage cost:  $V-1$  int ptrs +  $V(V-1)/2$  ints (but still  $O(V^2)$ )

Requires us to always use edges  $(v, w)$  such that  $v < w$ .

# Adjacency List Representation


43/110

For each vertex, store linked list of adjacent vertices:



Undirected graph

$A[0] = \langle 1, 3 \rangle$   
 $A[1] = \langle 0, 3 \rangle$   
 $A[2] = \langle 3 \rangle$   
 $A[3] = \langle 0, 1, 2 \rangle$



Directed graph

$A[0] = \langle 3 \rangle$   
 $A[1] = \langle 0, 3 \rangle$   
 $A[2] = \langle \rangle$   
 $A[3] = \langle 2 \rangle$

---

## ... Adjacency List Representation

44/110

### Advantages

- relatively easy to implement in languages like C
- can represent graphs and digraphs
- memory efficient if  $E:V$  relatively small

### Disadvantages:

- one graph has many possible representations  
(unless lists are ordered by same criterion e.g. ascending)

---

## ... Adjacency List Representation

45/110

### Graph initialisation

```
newGraph(V):  
    Input  number of nodes V  
    Output new empty graph  
  
    g.nV = V    // #vertices (numbered 0..V-1)  
    g.nE = 0    // #edges  
    allocate memory for g.edges[]  
    for all i=0..V-1 do  
        g.edges[i]=NULL    // empty list  
    end for  
    return g
```

---

## ... Adjacency List Representation

46/110

### Edge insertion:

```

insertEdge(g, (v,w)):
|   Input graph g, edge (v,w)
|
|   insertLL(g.edges[v], w)
|   insertLL(g.edges[w], v)
|   g.nE=g.nE+1

```

---

## ... Adjacency List Representation

47/110

Edge removal:

```

removeEdge(g, (v,w)):
|   Input graph g, edge (v,w)
|
|   deleteLL(g.edges[v], w)
|   deleteLL(g.edges[w], v)
|   g.nE=g.nE-1

```

---

## Exercise #6:

48/110

Analyse storage cost and time complexity of adjacency list representation

---

Storage cost:  $O(V+E)$  ( $V$  list pointers, total of  $2 \cdot E$  list elements)

- the larger of  $V, E$  determines the complexity

Cost of operations:

- initialisation:  $O(V)$  (initialise  $V$  lists)
- insert edge:  $O(1)$  (insert one vertex into list)
  - if you don't check for duplicates
- find/delete edge:  $O(V)$  (need to find vertex in list)

If vertex lists are sorted

- insert requires search of list  $\Rightarrow O(V)$
  - delete always requires a search, regardless of list order
- 

## Comparison of Graph Representations

50/110

	array of edges	adjacency matrix	adjacency list
space usage	$E$	$V^2$	$V+E$
initialise	$1$	$V^2$	$V$
insert edge	$1$	$1$	$1$
find/delete edge	$E$	$1$	$V$

Other operations:

	array of edges	adjacency matrix	adjacency list
disconnected(v)?	$E$	$V$	$1$
isPath(x,y)?	$E \cdot \log V$	$V^2$	$V+E$
copy graph	$E$	$V^2$	$E$
destroy graph	$1$	$V$	$E$

---

## Graph Abstract Data Type

### Graph ADT

52/110

Data:

- set of edges, set of vertices

Operations:

- building: create graph, add edge
- deleting: remove edge, drop whole graph
- scanning: check if graph contains a given edge

Things to note:

- set of vertices is fixed when graph initialised
- we treat vertices as `ints`, but could be arbitrary `Items`

---

### ... Graph ADT

53/110

Graph ADT interface `graph.h`

```
// graph representation is hidden
typedef struct GraphRep *Graph;

// vertices denoted by integers 0..N-1
typedef int Vertex;

// edges are pairs of vertices (end-points)
typedef struct Edge { Vertex v; Vertex w; } Edge;

// operations on graphs
Graph newGraph(int V);           // new graph with V vertices
int  numVertices(Graph);         // get number of vertices in a graph
void insertEdge(Graph, Edge);
void removeEdge(Graph, Edge);
bool adjacent(Graph, Vertex, Vertex); /* is there an edge
```

```
void showGraph(Graph);           between two vertices */
void freeGraph(Graph);          // print all edges in a graph
```


---

## Exercise #7: Graph ADT Client

54/110

Write a program that uses the graph ADT to

- build the graph depicted below
- print all the nodes that are incident to vertex 1 in ascending order



```
#include <stdio.h>
#include "Graph.h"

#define NODES 4
#define NODE_OF_INTEREST 1

int main(void) {
    Graph g = newGraph(NODES);

    Edge e;
    e.v = 0; e.w = 1; insertEdge(g, e);
    e.v = 0; e.w = 3; insertEdge(g, e);
    e.v = 1; e.w = 3; insertEdge(g, e);
    e.v = 3; e.w = 2; insertEdge(g, e);

    int v;
    for (v = 0; v < NODES; v++) {
        if (adjacent(g, v, NODE_OF_INTEREST))
            printf("%d\n", v);
    }

    freeGraph(g);
    return 0;
}
```



---

## Graph ADT (Array of Edges)

56/110

Implementation of GraphRep (array-of-edges representation)

```
typedef struct GraphRep {
    Edge *edges; // array of edges
    int nV;      // #vertices (numbered 0..nV-1)
    int nE;      // #edges
    int n;       // size of edge array
} GraphRep;
```





## Graph ADT (Adjacency Matrix)

57/110

### Implementation of GraphRep (adjacency-matrix representation)

```
typedef struct GraphRep {
    int **edges; // adjacency matrix
    int  nV;     // #vertices
    int  nE;     // #edges
} GraphRep;
```



## ... Graph ADT (Adjacency Matrix)

58/110

### Implementation of graph initialisation (adjacency-matrix representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);
    g->nV = V;  g->nE = 0;

    // allocate memory for each row
    g->edges = malloc(V * sizeof(int *));    assert(g->edges != NULL);
    // allocate memory for each column and initialise with 0
    for (i = 0; i < V; i++) {
        g->edges[i] = calloc(V, sizeof(int)); assert(g->edges[i] != NULL);
    }
    return g;
}
```

standard library function `calloc(size_t nelems, size_t nbytes)`

- allocates a memory block of size `nelems*nbytes`
- and sets all bytes in that block to *zero*

## ... Graph ADT (Adjacency Matrix)

59/110

## Implementation of edge insertion/removal (adjacency-matrix representation)

```
// check if vertex is valid in a graph
bool validV(Graph g, Vertex v) {
    return (g != NULL && v >= 0 && v < g->nV);
}

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!g->edges[e.v][e.w]) { // edge e not in graph
        g->edges[e.v][e.w] = 1;
        g->edges[e.w][e.v] = 1;
        g->nE++;
    }
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (g->edges[e.v][e.w]) { // edge e in graph
        g->edges[e.v][e.w] = 0;
        g->edges[e.w][e.v] = 0;
        g->nE--;
    }
}
```

---

### Exercise #8: Checking Neighbours (i)

60/110

Assuming an adjacency-matrix representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

---

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x) && validV(g,y));

    return (g->edges[x][y] != 0);
}
```

---

## Graph ADT (Adjacency List)


62/110

Implementation of GraphRep (adjacency-list representation)

```
typedef struct GraphRep {
    Node **edges; // array of lists
    int    nV;    // #vertices
    int    nE;    // #edges
} GraphRep;

typedef struct Node {
    Vertex    v;
    struct Node *next;
} Node;
```





## Exercise #9: Checking Neighbours (ii)

63/110

Assuming an adjacency list representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g, x));

    return inLL(g->edges[x], y);
}
```

inLL() checks if linked list contains an element

## Problems on Graphs

66/110

## Problems on Graphs

What kind of problems do we want to solve on/via graphs?

- is the graph fully-connected?
- can we remove an edge and keep it fully-connected?
- which vertices are reachable from  $v$ ? (transitive closure)
- is there a cycle that passes through all vertices? (circuit)
- what is the cheapest cost path from  $v$  to  $w$ ?
- is there a tree that links all vertices? (spanning tree)
- what is the minimum spanning tree?
- what is the maximal flow through a graph?
- ...
- can a graph be drawn in a plane with no crossing edges? (planar graphs)
- are two graphs "equivalent"? (isomorphism)
- ...

## Graph Algorithms

67/110

In this course we examine algorithms for

- graph traversal (simple graphs)
- reachability (directed graphs)
- minimum spanning trees (weighted graphs)
- shortest path (weighted graphs)

- maximum flow (weighted graphs)

---

## Graph Traversal

---

### Finding a Path

69/110

Questions on paths:

- is there a path between two given vertices ( $src, dest$ )?
- what is the sequence of vertices from  $src$  to  $dest$ ?

Approach to solving problem:

- examine vertices adjacent to  $src$
- if any of them is  $dest$ , then done
- otherwise try vertices two edges from  $src$
- repeat looking further and further from  $src$

Two strategies for graph traversal/search: *depth-first*, *breadth-first*


- DFS follows one path to completion before considering others
- BFS "fans-out" from the starting vertex ("spreading" subgraph)

---

### ... Finding a Path

70/110

Comparison of BFS/DFS search for checking if there is a path from  $a$  to  $h$  ...



Both approaches ignore some edges by remembering previously visited vertices.

---

## Depth-first Search

71/110

Depth-first search can be described recursively as

`depthFirst(G, v):`

1. mark  $v$  as visited
2. for each  $(v, w) \in \text{edges}(G)$  do  
if  $w$  has not been visited then

depthFirst(w)

The recursion induces *backtracking*

---

## ... Depth-first Search

72/110

Recursive DFS path checking


```
hasPath(G, src, dest):  
    Input  graph G, vertices src, dest  
    Output true if there is a path from src to dest in G,  
           false otherwise  
  
    mark all vertices in G as unvisited  
    return dfsPathCheck(G, src, dest)  
  
dfsPathCheck(G, v, dest):  
    mark v as visited  
    if v=dest then          // found dest  
        return true  
    else  
        for all (v,w) ∈ edges(G) do  
            if w has not been visited then  
                return dfsPathCheck(G, w, dest) // found path via w to dest  
            end if  
        end for  
    end if  
    return false           // no path from v to dest
```

---

## Exercise #10: Depth-first Traversal (i)

73/110

Trace the execution of dfsPathCheck(G, 0, 5) on:



Consider neighbours in ascending order

---

Answer:

0 - 1 - 2 - 3 - 4 - 5

---

## ... Depth-first Search

75/110

Cost analysis:

- all vertices marked as unvisited, each vertex visited at most once  $\Rightarrow$  cost =  $O(V)$
- visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(E)$ 
  - assuming an adjacency list representation

*Time complexity of DFS:  $O(V+E)$  (adjacency list representation)*

- the *larger* of  $V, E$  determines the complexity

For *dense graphs* ...  $E \cong V^2 \Rightarrow O(V+E) = O(V^2)$

For *sparse graphs* ...  $E \cong V \Rightarrow O(V+E) = O(V)$

## ... Depth-first Search

76/110

Note how different graph data structures affect cost:

- array-of-edges representation
  - visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(V \cdot E)$
  - cost of DFS:  $O(V \cdot E)$
- adjacency-matrix representation
  - visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(V^2)$
  - cost of DFS:  $O(V^2)$

## ... Depth-first Search

77/110

Knowing whether a path exists can be useful

Knowing what the path is even more useful

$\Rightarrow$  record the previously visited node as we search through the graph (so that we can then trace path through graph)

Make use of global variable:

- `visited[]` ... array to store previously visited node, for each node being visited

## ... Depth-first Search

78/110

`visited[]` // store previously visited node, for each vertex  $0..nV-1$

```
findPath(G, src, dest):
|   Input graph G, vertices src, dest
|
|   for all vertices v ∈ G do
|       visited[v] = -1
|   end for
|   visited[src] = src // starting node of the path
|   if dfsPathCheck(G, src, dest) then // show path in dest..src order
|       v = dest
|       while v ≠ src do
|           print v "-"
|           v = visited[v]
```

```

|   |   end while
|   |   print src
|   end if

dfsPathCheck(G, v, dest):
|   if v=dest then                // found edge from v to dest
|       return true
|   else
|       for all (v,w) ∈ edges(G) do
|           if visited[w]=-1 then
|               visited[w]=v
|               if dfsPathCheck(G, w, dest) then
|                   return true    // found path via w to dest
|               end if
|           end if
|       end for
|   end if
|   return false                  // no path from v to dest


```

---

## Exercise #11: Depth-first Traversal (ii)

79/110

Show the DFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	3	5	3	1	5	4	7	8
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-1-0

## ... Depth-first Search

81/110

DFS can also be described non-recursively (via a *stack*):

```

hasPath(G, src, dest):
|   Input  graph G, vertices src, dest
|   Output true if there is a path from src to dest in G,
|           false otherwise

```

```

mark all vertices in G as unvisited
push src onto new stack s
found=false
while not found and s is not empty do
    pop v from s
    mark v as visited
    if v=dest then
        found=true
    else
        for each (v,w) ∈ edges(G) such that w has not been visited
            push w onto s
        end for
    end if
end while
return found

```


Uses standard stack operations (push, pop, check if empty)

Time complexity is the same:  $O(V+E)$

## Exercise #12: Depth-first Traversal (iii)

82/110

Show how the stack evolves when executing `findPathDFS(g, 0, 5)` on:



Push neighbours in *descending* order ... so they get popped in ascending order

						4	5									
				3		5	5	5								
	1	2	4	4	4	4	4	4								
	4	4	4	4	4	4	4	4								
(empty)	→	0	→	5	→	5	→	5	→	5	→	5	→	5	→	5

## Breadth-first Search

84/110

Basic approach to breadth-first search (BFS):

- visit and mark current vertex
- visit all neighbours of current vertex
- then consider neighbours of neighbours

## Notes:

- tricky to describe recursively
- a minor variation on non-recursive DFS search works  
⇒ switch the *stack* for a *queue*

---

## ... Breadth-first Search

85/110

BFS algorithm (records visiting order, marks vertices as visited when put *on* queue):

visited[] // array of visiting orders, indexed by vertex 0..nV-1

```
findPathBFS(G, src, dest):  
    Input  graph G, vertices src, dest  
  
    for all vertices v ∈ G do  
        visited[v] = -1  
    end for  
    enqueue src into new queue q  
    visited[src] = src  
    found = false  
    while not found and q is not empty do  
        dequeue v from q  
        if v = dest then  
            found = true  
        else  
            for each (v, w) ∈ edges(G) such that visited[w] = -1 do  
                enqueue w into q  
                visited[w] = v  
            end for  
        end if  
    end while  
    if found then  
        display path in dest..src order  
    end if
```


Uses standard queue operations (enqueue, dequeue, check if empty)

---

## Exercise #13: Breadth-first Traversal

86/110

Show the BFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	0	2	5	0	5	5	3	-1
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-0

## ... Breadth-first Search

88/110

*Time complexity of BFS:  $O(V+E)$*  (adjacency list representation, same as DFS)

BFS finds a "shortest" path

- based on minimum # edges between *src* and *dest*.
- stops with first-found path, if there are multiple ones

In many applications, edges are weighted and we want path

- based on minimum sum-of-weights along path *src* .. *dest*

We discuss weighted/directed graphs later.

## Other DFS Examples

89/110

Other problems to solve via DFS graph search

- checking for the existence of a cycle
- determining which connected component each vertex is in



Graph with two connected components, a *path* and a *cycle*

## Exercise #14: Buggy Cycle Check

90/110

A graph has a *cycle* if

- it has a path of length  $> 1$
- with start vertex *src* = end vertex *dest*
- and without using any edge more than once

We are not required to give the path, just indicate its presence.

The following DFS cycle check has two bugs. Find them.



```

hasCycle(G):
|   Input  graph G
|   Output true if G has a cycle, false otherwise
|
|   choose any vertex  $v \in G$ 
|   return dfsCycleCheck(G, v)

dfsCycleCheck(G, v):
|   mark v as visited
|   for each  $(v, w) \in \text{edges}(G)$  do
|       if w has been visited then    // found cycle
|           return true
|       else if dfsCycleCheck(G, w) then
|           return true
|   end for
|   return false                      // no cycle at v

```

---

1. Only one connected component is checked.
2. The loop

```

    for each  $(v, w) \in \text{edges}(G)$  do

```

should exclude the neighbour of  $v$  from which you just came, so as to prevent a single edge  $w-v$  from being classified as a cycle.

---

## Computing Connected Components

92/110

Problems:

- how many connected subgraphs are there?
- are two vertices in the same connected subgraph?

Both of the above can be solved if we can

- build an array, one element for each vertex  $V$
  - indicating which connected component  $V$  is in
  - `componentOf[]` ... array  $[0..nV-1]$  of component IDs
- 

## ... Computing Connected Components

93/110

Algorithm to assign vertices to connected components:

```

components(G):
|   Input  graph G
|
|   for all vertices  $v \in G$  do
|       componentOf[v] = -1
|   end for
|   compID = 0
|   for all vertices  $v \in G$  do
|       if componentOf[v] = -1 then

```

```

| |         dfsComponents(G, v, compID)
| |         compID=compID+1
| |     end if
| end for

dfsComponents(G, v, id):
| componentOf[v]=id
| for all vertices w adjacent to v do
|     if componentOf[w]=-1 then
|         dfsComponents(G, w, id)
|     end if
| end for


```

## Exercise #15: Connected components

94/110

Trace the execution of the algorithm

1. on the graph shown below
2. on the same graph but with the dotted edges added



Consider neighbours in ascending order

1.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	-1	0	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
0	0	0	1	-1	-1	-1	-1
...							
0	0	0	1	1	2	2	2

2.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	0	-1	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
...							

---

## Hamiltonian and Euler Paths

---

### Hamiltonian Path and Circuit

97/110

*Hamiltonian path* problem:

- find a path connecting two vertices  $v, w$  in graph  $G$
- such that the path includes each *vertex* exactly once

If  $v = w$ , then we have a *Hamiltonian circuit*

Simple to state, but difficult to solve (*NP*-complete)

Many real-world applications require you to visit all vertices of a graph:

- Travelling salesman
- Bus routes
- ...


Problem named after Irish mathematician, physicist and astronomer Sir William Rowan Hamilton (1805 — 1865)

---

### ... Hamiltonian Path and Circuit

98/110

Graph and two possible Hamiltonian paths:



---

### ... Hamiltonian Path and Circuit

99/110

Approach:

- generate all possible simple paths (using e.g. DFS)
- keep a counter of vertices visited in current path
- stop when find a path containing  $V$  vertices

Can be expressed via a recursive DFS algorithm

- similar to simple path finding approach, except
  - keeps track of path length; succeeds if length =  $v-1$  (length =  $v$  for circuit)
  - resets "visited" marker after unsuccessful path

Algorithm for finding Hamiltonian path:

visited[] // array [0..nV-1] to keep track of visited vertices


```
hasHamiltonianPath(G, src, dest):
|   for all vertices v ∈ G do
|       visited[v] = false
|   end for
|   return hamiltonR(G, src, dest, #vertices(G)-1)
```

```
hamiltonR(G, v, dest, d):
|   Input G      graph
|       v      current vertex considered
|       dest destination vertex
|       d      distance "remaining" until path found
|
|   if v = dest then
|       if d = 0 then return true else return false
|   else
|       mark v as visited
|       for each unvisited neighbour w of v in G do
|           if hamiltonR(G, w, dest, d-1) then
|               return true
|           end if
|       end for
|   end if
|   mark v as unvisited          // reset visited mark
|   return false
```

## Exercise #16: Hamiltonian Path

101/110

Trace the execution of the algorithm when searching for a Hamiltonian path from 1 to 6:



Consider neighbours in ascending order

1-0-2-3-4-5-6	d ≠ 0
1-0-2-3-4-5-7-8-9	no unvisited neighbour
1-0-2-3-4-5-7-9-8	no unvisited neighbour
1-0-2-3-4-7-5-6	d ≠ 0
1-0-2-3-4-7-8-9	no unvisited neighbour
1-0-2-3-4-7-9-8	no unvisited neighbour

1-0-2-3-4-8-7-5-6  
 1-0-2-3-4-8-7-9  
 1-0-2-3-4-8-9-7-5-6

$d \neq 0$   
 no unvisited neighbour  
 ✓


Repeat on your own with  $\text{src}=0$  and  $\text{dest}=6$

## ... Hamiltonian Path and Circuit

103/110

Analysis: worst case requires  $(V-1)!$  paths to be examined

Consider a graph with isolated vertex and the rest fully-connected



Checking  $\text{hasHamiltonianPath}(g, x, 0)$  for any  $x$

- requires us to consider every possible path
- e.g 1-2-3-4, 1-2-4-3, 1-3-2-4, 1-3-4-2, 1-4-2-3, ...
- starting from any  $x$ , there are  $3!$  paths  $\Rightarrow 4!$  total paths
- there is no path of length 5 in these  $(V-1)!$  possibilities

There is no known simpler algorithm for this task  $\Rightarrow$  *NP*-hard.

Note, however, that the above case could be solved in constant time if we had a fast check for 0 and  $x$  being in the same connected component


## Euler Path and Circuit

104/110


*Euler path* problem:

- find a path connecting two vertices  $v, w$  in graph  $G$
- such that the path includes each *edge* exactly once  
 (note: the path does not have to be simple  $\Rightarrow$  can visit vertices more than once)

If  $v = w$ , then we have an *Euler circuit*



Euler Path: 4-3-1-5-2-1



Euler Circuit: 1-2-5-4-3-1

Many real-world applications require you to visit all edges of a graph:

- Postman
- Garbage pickup
- ...

## ... Euler Path and Circuit

105/110

One possible "brute-force" approach:

- check for each path if it's an Euler path
- would result in factorial time performance

Can develop a better algorithm by exploiting:

*Theorem.* A graph has an Euler circuit if and only if it is connected and all vertices have even degree


*Theorem.* A graph has a non-circuitous Euler path if and only if it is connected and exactly two vertices have odd degree

---

## Exercise #17: Euler Paths and Circuits

106/110

Which of these two graphs have an Euler path? an Euler circuit?



---

No Euler circuit

Only the second graph has an Euler path, e.g. 2-0-1-3-2-4-5

---

## ... Euler Path and Circuit

108/110

Assume the existence of  $\text{degree}(g, v)$  (degree of a vertex, cf. homework exercise 2 this week)

Algorithm to check whether a graph has an Euler path:

```
hasEulerPath(G, src, dest):  
|   Input  graph G, vertices src, dest  
|   Output true if G has Euler path from src to dest  
|           false otherwise  
|
```

```

if src≠dest then           // non~circuitous path
  if degree(G,src) or degree(G,dest) is even then
    return false
  end if
else if degree(G,src) is odd then // circuit
  return false
end if
for all vertices v∈G do
  if v≠src and v≠dest and degree(G,v) is odd then
    return false
  end if
end for
return true

```

---

## ... Euler Path and Circuit

109/110

Analysis of hasEulerPath algorithm:

- assume that connectivity is already checked
- assume that degree is available via  $O(1)$  lookup
- single loop over all vertices  $\Rightarrow O(V)$

If degree requires iteration over vertices

- cost to compute degree of a single vertex is  $O(V)$
- overall cost is  $O(V^2)$

$\Rightarrow$  problem tractable, even for large graphs (unlike Hamiltonian path problem)

For the keen: Linear-time (in the number of edges,  $E$ ) algorithm to compute an Euler path described in [Sedgewick] Ch.17.7

---

## Summary

110/110

- Graph terminology
  - vertices, edges, vertex degree, connected graph, tree
  - path, cycle, clique, spanning tree, spanning forest
- Graph representations
  - array of edges
  - adjacency matrix
  - adjacency lists
- Graph traversal
  - depth-first search (DFS)
  - breadth-first search (BFS)
  - cycle check, connected components
  - Hamiltonian paths/circuits, Euler paths/circuits
- Suggested reading (Sedgewick):
  - graph representations ... Ch. 17.1-17.5

- Hamiltonian/Euler paths ... Ch. 17.7
  - graph search ... Ch. 18.1-18.3, 18.7
- 

Produced: 3 Oct 2022