

## Tree Review

6/87

*Binary search trees ...*

- data structures designed for  $O(\log n)$  search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering ( $\text{data}(\text{Left}) < \text{root} < \text{data}(\text{Right})$ )
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (*worst:  $O(n)$* )
- operations: insert, delete, search, ...


## Balanced Search Trees

### Balanced BSTs

8/87

Reminder ...

- Goal: build binary search trees which have
  - minimum height  $\Rightarrow$  minimum worst case search cost
- Best balance you can achieve for tree with  $N$  nodes:
  - tree height of  $\log_2 N \Rightarrow$  worst case search  $O(\log N)$



Three *strategies* to improving worst case search in BSTs:

- *randomise* — reduce chance of worst-case scenario occurring
- *amortise* — do more work at insertion to make search faster
- *optimise* — implement all operations with performance bounds

9/87

# Rebalancing Trees

An approach to balanced trees:

- insert into leaves as for simple BST
- periodically, rebalance the tree

Question: how frequently/when/how to rebalance?

```
NewTreeInsert(tree, item):  
|   Input  tree, item  
|   Output tree with item randomly inserted  
|  
|   t=insertAtLeaf(tree, item)  
|   if #nodes(t) mod k = 0 then  
|       t=rebalance(t)  
|   end if  
|   return t
```

E.g. rebalance after every 20 insertions  $\Rightarrow$  choose  $k=20$

Note: To do this efficiently we would need to change tree data structure and basic operations:


```
typedef struct Node {  
    Item data;  
    int  nnodes;    // #nodes in my tree  
    Tree left, right; // subtrees  
} Node;
```

---

## ... Rebalancing Trees

10/87

How to rebalance a BST? Move median item to root.



---

## ... Rebalancing Trees

11/87

Implementation of rebalance:

```
rebalance(t):  
|   Input  tree t with n nodes  
|   Output t rebalanced  
|  
|   if  $n \geq 3$  then  
|       |   t=partition(t,  $\lfloor n/2 \rfloor$ )    // put node with median key at root  
|       |   left(t)=rebalance(left(t))    // then rebalance each subtree
```

```

|   |   right(t)=rebalance(right(t))
|   end if
|   return t

```


---

## ... Rebalancing Trees

12/87

New operation on trees:

- `partition(tree, i)`: re-arrange tree so that element with index  $i$  becomes root




For tree with  $N$  nodes, indices are  $0 \dots N-1$

---

## ... Rebalancing Trees

13/87

Partition: moves  $i^{\text{th}}$  node to root



## ... Rebalancing Trees

14/87

Implementation of partition operation:

```

partition(tree, i):
|   Input  tree with n nodes, index i
|   Output tree with item #i moved to the root
|
|   m=#nodes(left(tree))
|   if i < m then
|       left(tree)=partition(left(tree), i)
|       tree=rotateRight(tree)
|   else if i > m then
|       right(tree)=partition(right(tree), i-m-1)
|       tree=rotateLeft(tree)
|   end if
|   return tree


```

Note:  $\text{size}(\text{tree}) = n$ ,  $\text{size}(\text{left}(\text{tree})) = m$ ,  $\text{size}(\text{right}(\text{tree})) = n-m-1$  (why -1?)


## Exercise #1: Partition

15/87

Consider the tree  $t$ :



Show the result of  $\text{partition}(t, 3)$



## ... Rebalancing Trees

17/87

Analysis of rebalancing: visits every node  $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every  $k$  insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely  $\Rightarrow$  Solution: real balanced trees (later)

## Splay Trees

## Splay Trees

19/87

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering *parent-child-grandchild* (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.

---

## ... Splay Trees

20/87

Splay tree implementations also do *rotation-in-search*:

- by performing double-rotations also when searching

The idea: provides similar effect to periodic rebalance.

⇒ improves balance but makes search more expensive


---

## ... Splay Trees

21/87

Cases for splay tree double-rotations:


- case 1: grandchild is left-child of left-child ⇒ double right rotation from top
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child ⇒ double left rotation from top



## ... Splay Trees

22/87

Double-rotation case for left-child of left-child ("zig-zig"):




Note: both rotations at the root (unlike insertion-at-root)

---

## ... Splay Trees

23/87

Double-rotation case for right-child of left-child ("zig-zag"):



Note: rotate subtree first (like insertion-at-root)

## ... Splay Trees

24/87

Algorithm for splay tree insertion:



```
insertSplay(tree, item):
    Input  tree, item
    Output tree with item splay-inserted

    if tree is empty then return new node containing item
    else if item=data(tree) then return tree
    else if item<data(tree) then
        if left(tree) is empty then
            left(tree)=new node containing item
        else if item<data(left(tree)) then
            // Case 1: left-child of left-child "zig-zig"
            left(left(tree))=insertSplay(left(left(tree)), item)
            tree=rotateRight(tree)
        else if item>data(left(tree)) then
            // Case 2: right-child of left-child "zig-zag"
            right(left(tree))=insertSplay(right(left(tree)), item)
            left(tree)=rotateLeft(left(tree))
        end if
        return rotateRight(tree)
    else // item>data(tree)
        if right(tree) is empty then
            right(tree)=new node containing item
        else if item<data(right(tree)) then
            // Case 3: left-child of right-child "zag-zig"
            left(right(tree))=insertSplay(left(right(tree)), item)
            right(tree)=rotateRight(right(tree))
        else if item>data(right(tree)) then
            // Case 4: right-child of right-child "zag-zag"
            right(right(tree))=insertSplay(right(right(tree)), item)
            tree=rotateLeft(tree)
        end if
        return rotateLeft(tree)
    end if
```

## Exercise #2: Splay Trees

25/87

Insert 18 into this splay tree:



## ... Splay Trees

27/87

Searching in splay trees:

```


searchSplay(tree, item):
|   Input  tree, item
|   Output address of item if found in tree
|   NULL otherwise
|
|   if tree=NULL then
|       return NULL
|   else
|       tree=splay(tree, item)
|       if data(tree)=item then
|           return tree
|       else
|           return NULL
|       end if
|   end if
  
```

where `splay()` is similar to `insertSplay()`,  
except that it doesn't add a node ... simply moves `item` to root if found, or nearest node if not found


## Exercise #3: Splay Trees

28/87

If we search for 22 in the splay tree



... how does this affect the tree?




---

## ... Splay Trees

30/87

Why take into account both child and grandchild?

- moves accessed node to the root
- *moves every ancestor of accessed node roughly halfway to the root*

⇒ better amortized cost than insert-at-root

---

## ... Splay Trees

31/87

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider:  $m$  insert+search operations,  $n$  nodes
- *Theorem.* Total number of comparisons: average  $O((n+m) \cdot \log(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
  - improves balance on each search
  - moves frequently accessed nodes closer to root

But ... still has worst-case search cost  $O(n)$

---

## Real Balanced Trees

---

### Better Balanced Binary Search Trees

33/87

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have  $O(n)$  worst case

Ideally, we want both average/worst case to be  $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance



- red-black trees ... isomorphic to 2-3-4, but binary nodes

---

## AVL Trees

---

## AVL Trees

35/87

Invented by Georgy [Adelson-Velsky](#) and Evgenii [Landis](#)

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when:  $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by at most two rotations:

- if left subtree too deep ...
  - if data inserted in left-right grandchild  $\Rightarrow$  left-rotate left subtree
  - rotate right
- if right subtree too deep ...
  - if data inserted in right-left grandchild  $\Rightarrow$  right-rotate right subtree
  - rotate left

Problem: determining height/depth of subtrees may be expensive.

---

## ... AVL Trees

36/87

Implementation of AVL insertion

```
insertAVL(tree, item):
| Input  tree, item
| Output tree with item AVL-inserted
|
| if tree is empty then
|   return new node containing item
| else if item=data(tree) then
|   return tree
| else
|   if item<data(tree) then
|     left(tree)=insertAVL(left(tree), item)
|   else if item>data(tree) then
|     right(tree)=insertAVL(right(tree), item)
|   end if
|   if height(left(tree))-height(right(tree)) > 1 then
|     if item>data(left(tree)) then
|       left(tree)=rotateLeft(left(tree))
|     end if
|     tree=rotateRight(tree)
|   else if height(right(tree))-height(left(tree)) > 1 then
|     if item<data(right(tree)) then
|       right(tree)=rotateRight(right(tree))
|     end if
```


```
| | tree=rotateLeft(tree)
| | end if
| | return tree
| end if
```

---

## Exercise #4: AVL Trees

37/87

Insert 27 into the AVL tree



What would happen if you now insert 28?

You may like the animation at [www.cs.usfca.edu/~galles/visualization/AVLtree.html](http://www.cs.usfca.edu/~galles/visualization/AVLtree.html)

---

## ... AVL Trees

39/87

Analysis of AVL trees:

- trees are *height*-balanced; subtree depths differ by  $\pm 1$
- average/worst-case search performance of  $O(\log n)$
- *require* extra data to be stored in each node ("height")
- may not be *weight*-balanced; subtree sizes may differ




## 2-3-4 Trees

## 2-3-4 Trees

41/87

2-3-4 trees have three kinds of nodes

- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



## ... 2-3-4 Trees

42/87

2-3-4 trees are ordered similarly to BSTs



In a *balanced 2-3-4 tree*:

- all leaves are at same distance from the root

2-3-4 trees grow "upwards" by splitting 4-nodes.

## ... 2-3-4 Trees

43/87

Possible 2-3-4 tree data structure:

```

typedef struct node {
    int         order;    // 2, 3 or 4
    int         data[3];  // items in node
    struct node *child[4]; // links to subtrees
} node;

```

## ... 2-3-4 Trees

44/87

Searching in 2-3-4 trees:

**Search(tree, item):**

```

Input tree, item
Output address of item if found in 2-3-4 tree
NULL otherwise

if tree is empty then
    return NULL
else
    i=0
    while i<tree.order-1 and item>tree.data[i] do
        i=i+1 // find relevant slot in data[]
    end while
    if item=tree.data[i] then // item found
        return address of tree.data[i]
    else // keep looking in relevant subtree

```

```

| |         return Search(tree.child[i], item)
| |     end if
| end if

```

---

## ... 2-3-4 Trees

45/87

2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height  $h$
  - 2-3-4 trees are always balanced  $\Rightarrow$  height is  $O(\log n)$
  - worst case for height: all nodes are 2-nodes  
same case as for balanced BSTs, i.e.  $h \cong \log_2 n$
  - best case for height: all nodes are 4-nodes  
balanced tree with branching factor 4, i.e.  $h \cong \log_4 n$
- 

## Insertion into 2-3-4 Trees

46/87

Starting with the root node:

**repeat**

- if current node is full (i.e. contains 3 items)
  - split into two 2-nodes
  - promote middle element to parent
    - if no parent  $\Rightarrow$  middle element becomes the new root 2-node
  - go back to parent node
- if current node is a leaf
  - insert Item in this node, order++
- if current node is not a leaf
  - go to child where Item belongs


**until** Item inserted

---


## ... Insertion into 2-3-4 Trees

47/87

Insertion into a 2-node or 3-node:




Insertion into a 4-node (requires a split):



## ... Insertion into 2-3-4 Trees

48/87


Splitting the root:



## ... Insertion into 2-3-4 Trees

49/87



Building a 2-3-4 tree ... 7 insertions:



## Exercise #5: Insertion into 2-3-4 Tree

50/87

Show what happens when D, S, F, U are inserted into this tree:



## ... Insertion into 2-3-4 Trees

52/87

Insertion algorithm:

```

insert(tree, item):
|   Input  2-3-4 tree, item
|   Output tree with item inserted
|
|   node=root(tree), parent=NULL
|   repeat
|       if node.order=4 then
|           promote = node.data[1]      // middle value
|           nodeL   = new node containing node.data[0]
|           nodeR   = new node containing node.data[2]
|           if parent=NULL then
|               make new 2-node root with promote,nodeL,nodeR
|           else
|               insert promote,nodeL,nodeR into parent
|               increment parent.order
|           end if
|           node=parent
|       end if
|       if node is a leaf then
|           insert item into node
|           increment node.order
|       else
|           parent=node
|           i=0
|           while i<node.order-1 and item>node.data[i] do
|               i=i+1      // find relevant child to insert item
|           end while
|           node=node.child[i]
|       end if
|   until item inserted

```

---

## ... Insertion into 2-3-4 Trees

53/87

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or  $M$ -way trees?

- allow nodes to hold up to  $M-1$  items, and at least  $M/2$
- if each node is a disk-page, then we have a *B-tree* (databases)
- for B-trees, depending on  $\text{Item size}$ ,  $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees  $\rightarrow$  red-black trees.

---

## Red-Black Trees

## Red-Black Trees

55/87

*Red-black trees* are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- *red* links ... combine nodes to represent 3- and 4-nodes
- *black* links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

## Red-Black Trees

56/87

Definition of a *red-black tree*

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 *sibling* of its parent
- a black node corresponds to a 2-3-4 *child* of its parent
  - if no parent (= root) → also black

*Balanced* red-black tree

- all paths from root to leaf have same number of black nodes


Insertion algorithm: avoids worst case  $O(n)$  behaviour

Search algorithm: standard BST search

## ... Red-Black Trees

57/87

Representing 4-nodes in red-black trees:




Some texts colour the links rather than the nodes.

## ... Red-Black Trees

58/87

Representing 3-nodes in red-black trees (two possibilities):



## ... Red-Black Trees

59/87

Equivalent trees (one 2-3-4, one red-black):



## ... Red-Black Trees

60/87

Red-black tree implementation:

```
typedef enum {RED, BLACK} Colr;
typedef struct node *RBTree;
typedef struct node {
    Item data; // actual data
    Colr color; // relationship to parent
    RBTree left; // left subtree
    RBTree right; // right subtree
} node;

#define color(tree) ((tree)->color)
#define NodeisRed(t) ((t) != NULL && (t)->color == RED)

RED = node is part of the same 2-3-4 node as its parent (sibling)

BLACK = node is a child of the 2-3-4 node containing the parent
```

## ... Red-Black Trees

61/87

New nodes are always **red**:



```

RBTree newNode(Item it) {
    RBTree new = malloc(sizeof(Node));
    assert(new != NULL);
    data(new) = it;
    color(new) = RED;
    left(new) = right(new) = NULL;
    return new;
}

```


---

## ... Red-Black Trees

62/87

Node.color allows us to distinguish links

- black = parent node is a "real" parent
- red = parent node is a 2-3-4 neighbour



## ... Red-Black Trees

63/87

Search method is standard BST search:

```

SearchRedBlack(tree, item):
    Input  tree, item
    Output true if item found in red-black tree
           false otherwise

    if tree is empty then
        return false
    else if item < data(tree) then
        return SearchRedBlack(left(tree), item)
    else if item > data(tree) then
        return SearchRedBlack(right(tree), item)
    else // found
        return true
    end if

```

---

## Red-Black Tree Insertion

64/87

Insertion is more complex than for standard BSTs

- need to recall direction of last branch (left or right)
  - splitting/promoting implemented by rotateLeft/rotateRight
  - several cases to consider depending on colour/direction combinations
- 

65/87

## ... Red-Black Tree Insertion

High-level description of insertion algorithm:

`insertRB(tree, item, inRight):`

```
| Input tree, item, inRight indicating direction of last branch
| Output tree with it inserted
|
| if tree is empty then
|     return newNode(item)
| else if item=data(tree) then
|     return tree
| end if
| if left(tree) and right(tree) both are RED then
|     split 4-node in a red-black tree
| end if
| recursive insert a la BST, re-arrange links/colours after insert
| return modified tree
```


`insertRedBlack(tree, item):`

```
| Input red-black tree, item
| Output tree with item inserted
|
| tree=insertRB(tree, item, false)
| color(tree)=BLACK
| return tree
```

## ... Red-Black Tree Insertion

66/87

Splitting a 4-node, in a red-black tree:



Algorithm:

```
color(left(currentTree))=BLACK
color(right(currentTree))=BLACK
color(currentTree)=RED
```

## ... Red-Black Tree Insertion

67/87

Simple recursive insert (a la BST):



Algorithm:

```

if item < data(tree) then
    left(tree) = insertRB(left(tree), item, false)
    re-arrange links/colours after insert
else // item larger than data in root
    right(tree) = insertRB(right(tree), item, true)
    re-arrange links/colours after insert
end if

```

Not affected by colour of tree node.

## ... Red-Black Tree Insertion

68/87

Re-arrange links/colours after insert:

Step 1 — "normalise" direction of successive red links



Algorithm:

```

if inRight and both t and left(t) are red then
    t = rotateRight(t)
end if

```

Symmetrically,


- if not inRight and both t and right(t) are red  
⇒ left rotate current tree t

## ... Red-Black Tree Insertion

69/87

Re-arrange links/colours after insert:

Step 2 — two successive red links = newly-created 4-node



## Algorithm:

```
if both left child and left-left grandchild are red then
    color(t)=RED
    color(left(t))=BLACK
    t=rotateRight(t)
end if
```

Symmetrically,

- if both right child and right-right grandchild are red  
⇒ re-colour current tree  $t$  and  $\text{right}(t)$ , then left rotate  $t$

---

## ... Red-Black Tree Insertion

70/87

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



---

## Red-black Tree Performance

71/87

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is  $O(\log_2 n)$
- insertion affects nodes down one path; #rotations+recolourings is  $O(h)$   
(where  $h$  is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgwick.

---

## Tries

### ... Tries

73/87


A *trie* ...

- is a compact data structure for representing a set of strings
  - e.g. all the words in a text, a dictionary etc.

Note: Trie comes from *retrieval*, but is pronounced like "try" to distinguish it from "tree"

---

*Tries* are trees organised using parts of keys (rather than whole keys)



## Exercise #6:

75/87

How many words are encoded in the trie on the previous slide?

11

## ... Tries

77/87

Each node in a trie ...

- contains one part of a key (typically one character)
- may have up to 26 children
- may be tagged as a "finishing" node
- but even "finishing" nodes may have children

*Depth*  $d$  of trie = length of longest key value

Cost of searching  $O(d)$  (independent of  $n$ )

## ... Tries

78/87

Possible trie representation:

```
#define ALPHABET_SIZE 26

typedef struct Node *Trie;


typedef struct Node {
    bool finish;        // last char in key?
    Item data;          // no Item if !finish
    Trie child[ALPHABET_SIZE];
} Node;
```

```
typedef char *Key;
```

## ... Tries

79/87

Note: Can also use BST-like nodes for more space-efficient implementation of tries



## Trie Operations

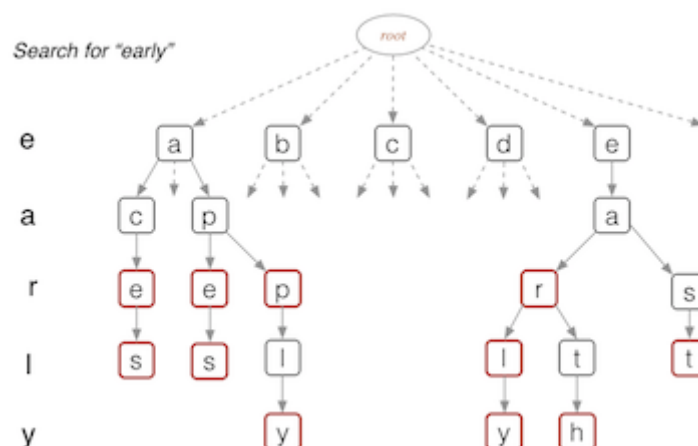
80/87

Basic operations on tries:

1. search for a key
2. insert a key

## ... Trie Operations

81/87



## ... Trie Operations

82/87

Traversing a path, using char-by-char from Key:

```
find(trie, key):
| Input  trie, key
| Output pointer to element in trie if key found
|         NULL otherwise
```

```

node=trie
for each char in key do
|   if node.child[char] exists then
|       node=node.child[char]  // move down one level
|   else
|       return NULL
|   end if
end for
if node.finish then          // "finishing" node reached?
    return node
else
    return NULL
end if

```

---

## ... Trie Operations

83/87

### Insertion into Trie:

```

insert(trie, item, key):
|   Input  trie, item with key of length m
|   Output trie with item inserted
|
|   if trie is empty then
|       t=new trie node
|   end if
|   if m=0 then
|       t.finish=true, t.data=item
|   else
|       t.child[key[0]]=insert(t.child[key[0]], item, key[1..m-1])
|   end if
|   return t


```


---

## Exercise #7: Trie Insertion

84/87

Insert **cat**, **cats** and **carer** into this trie:





## ... Trie Operations

86/87

Analysis of standard tries:

- $O(n)$  space
- insertion and search in time  $O(m)$ 
  - $n$  ... total size of text (e.g. sum of lengths of all strings in a given dictionary)
  - $m$  ... size of the string parameter of the operation (the "key")

## Summary

87/87

- Tree operations
  - tree partition
  - rebalancing
- Self-adjusting trees
  - Splay trees
  - AVL trees
  - 2-3-4 trees
  - Red-black trees
- Tries
- Suggested reading:
  - Sedgewick, Ch. 12.9
  - Sedgewick, Ch. 13.1-13.4
  - Sedgewick, Ch. 15.2